

GhostWorker AI

Autonomous Workplace Digital Twins Living in Slack, GitHub & Meeting Memory — Eliminating the Async Knowledge Bottleneck.



Core Innovation

Leaving the chatbox: An ambient AI digital twin of every key team member that automatically steps in to answer blockers with verifiable citations.



Seamless Integration

Operates directly inside Slack threads, GitHub RFCs, Jira issues, and voice briefings without demanding a new portal or workflow.



Edge & On-Device Ready

Designed for hybrid edge-cloud deployment: local quantized SLMs on device for personal memory, cloud neural grounding for deep synthesis.

THE PAIN POINT

The 24-Hour Knowledge Black Hole

Modern engineering and product teams waste up to 30% of their sprints waiting for asynchronous answers across time zones and meetings.



The Vacation & Meeting Stall

When a Tech Lead is asleep or stuck in 6 hours of back-to-back calls, a 30-second architectural query stalls the entire engineering pod for an entire business day.



The Generic Chatbot Failure

ChatGPT or Claude have zero access to private PR debates, unrecorded Slack agreements, or historical trade-offs. Engineers refuse to manually copy-paste internal secrets.



The "New Tab" Fatigue

Workplace tools fail when they require workers to adopt yet another website. True enterprise productivity happens natively inside existing team channels.

How GhostWorker Works Under the Hood

A three-tier multi-agent pipeline providing persona alignment, episodic memory, and verifiable citations.

Continuous Memory Engine

- ▶ **Episodic Ingestion:** Streams Slack discussions, Git commit diffs, and Jira tickets without manual training.
- ▶ **Fact Extraction & Weighting:** Continuously distills truths and calculates confidence scores as teammates communicate.
- ▶ **Deterministic Citations:** Every assertion generated is backed by a clickable link to the source Slack message or PR line.

Autonomous Multi-Agent Behaviors

- ▶ **In-Thread Ambient Interception:** Monitors unanswered threads; autonomously drafts verified technical unblockers.
- ▶ **Executive Multi-Clone Polling:** A CEO can ask *"Are we on track for v3?"* — polling Engineering, Sales, and Product twins in parallel to spot hidden blockers.
- ▶ **Voice & On-Device Edge:** Natural voice summaries via Whisper/TTS, with on-device local SLM inference for privacy-first notes.

Why Syed Squad Stands Out

Built by a seasoned enterprise engineer with proven delivery across modern AI agent architectures.



Proven Track Record

- ▶ Software Engineering Associate at Accenture (Healthcare & Life Sciences backend).
- ▶ Anthropic Claude Certified Architect.
- ▶ 31+ Cloud & AI Certifications (Google Cloud, Microsoft Azure).



Shipped Production Code

- ▶ Built & open-sourced Aegis MCP Registry, GhostWorker AI, and Enterprise RAG engines.
- ▶ Deep expertise in LLM tool calling, Model Context Protocol (MCP), and local model execution.



Hackathon Deliverables Ready

- ▶ Working functional prototype with simulated enterprise teams.
- ▶ Live public demonstration URL & open-source codebase.
- ▶ Clear roadmap for on-device Android smartphone optimization for iQOO devices.